

# Limitations of Ordinary Linear Regression

## 1. Multicollinearity

- Multicollinearity occurs when two or more independent variables in a regression model are highly correlated with each other.
- When independent variables are **highly correlated**, it becomes hard to estimate the coefficients accurately.
- Leads to **unstable** and **unreliable** models.
- Multicollinearity denotes when independent variables in a linear regression equation are correlated.



## Example: Height in Inches vs. Height in Meters

Person	Height (inches) $x_1$	Height (meters) $x_2$	Weight (kg) $y$
A	60	1.524	45
B	65	1.651	55
C	70	1.778	65
D	72	1.829	70
E	75	1.905	75



## Relationship

$$x_2 = x_1 \times 0.0254$$

This means **Height in meters is linearly dependent** on Height in inches → **perfect multicollinearity**.



## Why This Is a Problem in Regression

Suppose we try to model:

$$y = \beta_0 + \beta_1 \cdot \text{Height\_inches} + \beta_2 \cdot \text{Height\_meters} + \epsilon$$

The model will **fail to assign unique values** to  $\beta_1$  and  $\beta_2$ , since both inputs contain the **same information**.



## Correlation

$$\text{Corr}(x_1, x_2) \approx 1$$

Meaning they are **perfectly correlated**, causing:

- Inflated **standard errors**
- Unstable **coefficients**
- High **VIF values**



## 2. High Variance When Number of Predictors is Large

- When the number of **predictors (features/input/independent variables)** is close to or exceeds the number of **observations (samples)**, the model becomes **unstable**.
- This results in a model that:
  - Fits the **training data too well**
  - Performs **poorly on test/unseen data**
- This is a classic case of **overfitting**, where the model captures noise rather than true patterns.



## 3. Overfitting: Model Fits Noise Instead of Signal

- **Overfitting** happens when a model learns not only the **underlying patterns** in the data (signal), but also the **random fluctuations or errors** (noise).
- The model performs **very well on training data**, but **fails to generalize** to new or unseen data.



## Why It Happens:

- Too many predictors/features
- Very flexible models
- Small training datasets
- High variance in the data



### Example:

- Imagine fitting a **complex curve** to just a few data points. It may pass through all the points (perfect accuracy), but perform poorly on future data.



## 4. Poor Generalization to Unseen Data

- Generalization** refers to how well a model performs on **new, unseen data** — not just the data it was trained on.
- Ordinary Linear Regression** assumes:
  - A **linear relationship** between predictors and output
  - Constant variance** of errors (homoscedasticity)
  - Normally distributed** residuals
  - Independence** of observations



### When these assumptions are violated:

- The model may perform well on training data but **poorly on test data**.
- This leads to **low predictive accuracy** in real-world applications.



### Example:

If the true relationship is **non-linear**, a linear model will **underfit**, missing key patterns — resulting in **poor generalization**.

Limitation	Explanation	Example
<b>High Variance When Number of Predictors is Large</b>	Model becomes unstable when there are too many predictors compared to data size.	Predicting house prices with 10 features but only 8 houses in the dataset.
<b>Multicollinearity Causes Unstable Coefficient Estimates</b>	Highly correlated predictors distort the influence of individual variables.	Predicting salary using both <code>years of experience</code> and <code>age</code> , which are strongly correlated.
<b>Overfitting: Fits Noise Instead of Signal</b>	Model learns random noise rather than actual patterns in data.	A model gives perfect predictions on training data but fails completely on test data.

Limitation	Explanation	Example
<b>Poor Generalization to Unseen Data</b>	Model doesn't perform well on new or unseen data.	Linear regression used to predict non-linear relationships like stock prices or weather trends.



## 2. Polynomial Regression

### ✓ Concept

- Polynomial Regression is an extension of **Linear Regression** that allows for **non-linear** relationships between the **independent variable(s)** and the **dependent variable**.
- It does so by including **powers (exponents)** of the predictors as additional features.
- Despite modeling non-linear relationships, it is still considered a **linear model** in terms of the coefficients.



### Why Use It?

- When a scatterplot of the data shows a **curved trend** instead of a straight line.
- It helps capture **quadratic**, **cubic**, or even higher-order patterns in the data.



### Equation

For one predictor  $x$ , a **degree-2 (quadratic)** polynomial regression model looks like:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \varepsilon$$

More generally, for a degree- $d$  polynomial:

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \dots + \beta_d x^d + \varepsilon$$

Where:

- $y$ : Dependent variable (target)
- $x$ : Independent variable (predictor)
- $\beta_0, \beta_1, \dots, \beta_d$ : Coefficients
- $\varepsilon$ : Error term

## When to Use Polynomial Regression

### ✓ 1. Residual Plots Show Non-Random Patterns

- In a good linear model, residuals (errors) should be **randomly scattered** around zero.
- If the residual plot shows a **curve or systematic pattern**, it indicates that a **linear model is not appropriate**, and a **polynomial regression** may be needed.

### ✓ 2. $R^2$ Improves Significantly with Higher-Order Terms

- Add polynomial terms (e.g.,  $x^2$ ,  $x^3$ ) to the model and observe  $R^2$  (coefficient of determination).
- If  $R^2$  **increases noticeably** with added terms, it suggests the model is better capturing the underlying **non-linear relationship**.
- But beware: A very high degree may lead to **overfitting**.

### Example Workflow:

1. Fit a linear regression.
2. Plot the residuals.
  - If they curve → try a degree-2 polynomial.
3. Check  $R^2$ :
  - If  $R^2$  increases from, say, **0.65 to 0.91**, polynomial terms are likely adding real value.

## Regularization Techniques in Regression

Regularization helps prevent **overfitting** by adding a **penalty** to the model's complexity—specifically, to the **size of the coefficients**.

### 1. Ridge Regression (L2 Regularization)

- **Penalty:** Sum of squared coefficients

$$\text{Loss} = \text{RSS} + \lambda \sum_{j=1}^p \beta_j^2$$

- Shrinks all coefficients **toward zero** but **does not eliminate any**.
- **Use case:** When all predictors are important but you want to reduce model complexity.



## 2. Lasso Regression (L1 Regularization)

- **Penalty:** Sum of absolute coefficients

$$\text{Loss} = \text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$

- Can shrink **some coefficients exactly to zero**, effectively performing **feature selection**.
- **Use case:** When you expect **only a few variables** to be truly important.



## 3. Elastic Net Regression

- **Penalty:** Combination of L1 and L2

$$\text{Loss} = \text{RSS} + \lambda_1 \sum_{j=1}^p |\beta_j| + \lambda_2 \sum_{j=1}^p \beta_j^2$$

- Balances the **feature selection ability** of Lasso with the **stability** of Ridge.
- **Use case:** When predictors are highly correlated or you need both **shrinkage** and **selection**.






## Tuning the Regularization Parameter $\lambda$

- $\lambda$  controls the **strength of the penalty**:
  - **High  $\lambda$**  → More penalty → Coefficients are **shrunk more** → Model is **simpler**.
  - **Low  $\lambda$**  → Less penalty → Coefficients are closer to **OLS estimates**.
- Typically chosen using **cross-validation** (e.g., k-fold CV).



## Summary Table

Technique	Penalty Type	Effect on Coefficients	Feature Selection?
Ridge Regression	L2 (squared values)	Shrinks all coefficients, none set to 0	 No
Lasso Regression	L1 (absolute values)	Some coefficients shrunk to 0	 Yes
Elastic Net	L1 + L2 (combined penalty)	Shrinks and selects variables	 Yes